

PARALLEL AND SEQUENTIAL ALGORITHMS AND DATA STRUCTURES

LECTURE 11

Randomized Algorithms



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Randomized Algorithms

- Exploit randomness during computation
 - Pivot selection in Quicksort
 - Average case analysis
 - Primality testing
- **Question:** How many comparisons are needed to find the *second* largest number on a sequence of n numbers?
 - Naive algorithm: $2n-3$ comparisons **why?**
 - Divide and Conquer algorithm: $3n/2$ comparisons **why?**
 - Simple randomized algorithm: $n-1+2\log n$ comparisons *on the average* **why?**



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

SYNOPSIS

- Expectation versus High Probability
- Finding The Two Largest
- Find the k^{th} smallest element
- Quicksort
- Analysis of Quicksort



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Expected vs. High Probability Bounds.

- **Expected bounds**
 - the average case across all random choices used in the algorithm
 - Once in a while, the work could be much larger
- **High-probability bounds**
 - it is very unlikely that the cost will be above some bound
 - an algorithm on n elements has $O(n)$ work with probability at least $1-1/n^5$
 - This means that only once in about n^5 tries will the algorithm require more than $O(n)$ work



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Expected vs. High Probability Bounds.

- we had 100 students take exams, most of the time each student takes 1 hour, but that once on every 100 exams or so, each student gets hung up and takes 101 hours
 - The average for each student is $(99 \cdot 1 + 1 \cdot 101) / 100 = 2$ hours
 - the **expected maximum** will be close to 100 hours
 - ✓ on most exams with a hundred students one student will get hung up, so the expected maximum will be close to 100 hours, not 2 hours
 - Every student will finish in 2 hours with probability $1 - 1/n^5$



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Expected vs. High Probability Bounds.

- properties of summing vs. taking a maximum
 - analyze **work** using **expectation**
 - analyze **span** using **high probability**

why?



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

SYNOPSIS

- Discrete Probability: Let's Toss Some Dice
- Finding The Two Largest
- Find the k^{th} smallest element
- Quicksort
- Analysis of Quicksort



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Finding The Top Two Elements

```
1  max2 S =  
2  let  
3    replace ((m1, m2), v) =  
4      if v ≤ m2 then (m1, m2)  
5      else if v ≤ m1 then (m1, v)  
6      else (v, m1)  
7    val init = if S1 ≥ S2 then (S1, S2) else (S2, S1)  
8  in  
9    iter replace init S⟨3, ..., n⟩  
10 end
```

- We will do exact analysis
- $1+2(n-2)=2n-3$ comparisons in the worst case (Why?)
- A Divide and Conquer algorithm gives $3n/2-2$ (how?)



华中科技大学

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY



Worst Case Analysis



```
1  max2 S =  
2  let  
3    replace ((m1, m2), v) =  
4      if v ≤ m2 then (m1, m2)  
5      else if v ≤ m1 then (m1, v)  
6      else (v, m1)  
7    val init = if S1 ≥ S2 then (S1, S2) else (S2, S1)  
8  in  
9    iter replace init S⟨3, ..., n⟩  
10 end
```

- An already sorted sequence (e.g., $\langle 1, 2, 3, \dots, n \rangle$) will need exactly $2n-3$ comparisons
- But this happens with $1/n!$ chance (Why?)



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

A Randomized Algorithm

- The worst-case analysis is overly pessimistic (**why?**)
- Consider the following variant
 - On input of a sequence S of n elements
 - ✓ 1. Let $T = \text{permute}(S, \pi)$, where π is a random permutation (i.e., we choose one of the $n!$ permutations)
 - ✓ 2. Run the **max2** algorithm on T
 - No need to really generate the permutation!
 - ✓ Just pick an unprocessed element randomly until all elements are processed
 - ✓ It is convenient to model this by one initial permutation!



明

Analysis

```
1  max2 S =  
2  let  
3    replace ((m1, m2), v) =  
4      if v ≤ m2 then (m1, m2)  
5      else if v ≤ m1 then (m1, v)  
6      else (v, m1)  
7  val init = if S1 ≥ S2 then (S1, S2) else (S2, S1)  
8  in  
9    iter replace init S⟨3, ..., n⟩  
10 end
```

- X_i : an indicator random variable, denoting whether Line 5 gets executed for the value at S_i
- Y is the number of comparisons



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Analysis

```

1  max2 S =
2  let
3    replace ((m1, m2), v) =
4      if v ≤ m2 then (m1, m2)
5      else if v ≤ m1 then (m1, v)
6      else (v, m1)
7    val init = if S1 ≥ S2 then (S1, S2) else (S2, S1)
8  in
9    iter replace init S⟨3, ..., n⟩
10 end

```

• Y=?

$$Y(e) = \underbrace{1}_{\text{Line 7}} + \underbrace{n-2}_{\text{Line 4}} + \underbrace{\sum_{i=3}^n X_i(e)}_{\text{Line 5}}$$



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Analysis

- This expression is true regardless of the random choice we're making
- We're interested in computing the expected value of Y
- By linearity of expectation, $E[Y] = ?$

$$\begin{aligned} E[Y] &= E \left[1 + (n - 2) + \sum_{i=3}^n X_i \right] \\ &= 1 + (n - 2) + \sum_{i=3}^n E[X_i]. \end{aligned}$$



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Analysis

- Problem boils down to computing $E[X_i]$, for $i=3, \dots, n!$
- What is the probability that $T_i > m_2$?
 - $T_i > m_2$ holds when T_i is either the largest or the second largest in $\{T_1, \dots, T_i\}$
- So, what is the probability that T_i is one of the two largest elements in a randomly permuted sequence of length i ?
 - $1/i + 1/i = 2/i$
- $E[X_i] = 1 \cdot 2/i = 2/i$



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Analysis

德厚學
求是創新

$$\begin{aligned}\mathbf{E}[Y] &= 1 + (n - 2) + \sum_{i=3}^n \mathbf{E}[X_i] \\ &= 1 + (n - 2) + \sum_{i=3}^n \frac{2}{i} \\ &= 1 + (n - 2) + 2\left(\frac{1}{3} + \frac{1}{4} + \dots \frac{1}{n}\right) \\ &= n - 4 + 2\left(1 + \frac{1}{2} + \frac{1}{3} + \frac{1}{4} + \dots \frac{1}{n}\right) \\ &= n - 4 + 2H_n\end{aligned}$$

- H_n is the n^{th} Harmonic number
- $H_n \leq 1 + \log_2 n$
- $\mathbf{E}[Y] \leq n - 2 + 2\log_2 n$



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

SYNOPSIS

- Discrete Probability: Let's Toss Some Dice
- Finding The Two Largest
- Find the k^{th} smallest element
- Quicksort
- Analysis of Quicksort



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Finding The K^{th} Smallest Element

- **Input:** a sequence of n numbers (not necessarily sorted)
- **Output:** the K^{th} smallest value in S (i.e., (nth (sort S) k))
- **Requirement:** $O(n)$ expected work and $O(\log^2 n)$ span
- **We can't really sort the sequence!**



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Finding The K^{th} Smallest Element

```
1  kthSmallest(k, S) = let
2    p = S0
3    L =  $\langle x \in S \mid x < p \rangle$ 
4    R =  $\langle x \in S \mid x > p \rangle$ 
5  in
6    if (k < |L|) then kthSmallest(k, L)
7    else if (k < |S| - |R|) then p
8    else kthSmallest(k - (|S| - |R|), R)
```

- Let $X = \max\{|L|, |R|\} / |S|$, which

$$W(n) = W(X \cdot n) + O(n)$$

$$S(n) = S(X \cdot n) + O(\log n)$$

how?



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Finding The K^{th} Smallest Element


- We want to find $E[X]$?

➤ The probability that we land on a point on the x axis is $1/n$:

$$E[X] = \frac{1}{n} \sum_{i=0}^{n-1} \max\{i, n-i-1\}/n$$

➤ the size of L and size of R:

$$E[X] \leq \frac{1}{n} \sum_{j=n/2}^{n-1} \frac{2}{n} \cdot j \leq \frac{3}{4}$$



$$\sum_{i=x}^y i = \frac{1}{2}(x+y)(y-x+1).$$

Finding The K^{th} Smallest Element

- Theorem 10.16. Starting with size n , the expected size of S in algorithm *kthSmallest* after i recursive calls is $(3/4)^i n$
 - Let Y_i be the random variable representing the size of the result after step (recursive call) i

$$Y_i = n \prod_{j=1}^i X_j$$

$$\mathbf{E}[Y_i] = \mathbf{E}\left[n \prod_{j=1}^i X_j\right] = n \prod_{j=1}^i \mathbf{E}[X_j] \leq \left(\frac{3}{4}\right)^i n$$

why?

Finding The K^{th} Smallest Element

- $W=?$

➤ The work at each step is linear: $W_{contract}(n) \leq k_1 n + k_2$

$$\begin{aligned} \mathbf{E} [W_{kthSmallest}(n)] &\leq \sum_{i=0}^n (k_1 n \left(\frac{3}{4}\right)^i + k_2) \\ &\leq k_1 n \left(\sum_{i=0}^n \left(\frac{3}{4}\right)^i \right) + k_2 n \\ &\leq 4k_1 n + k_2 n \\ &\in O(n) \end{aligned}$$

why?



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Finding The K^{th} Smallest Element

- $S=?$
 - The span at each step is linear: $S_{contract}(n) \leq k_1 \log n + k_2$
- Bound the span of the algorithm
 - Consider step $i = 10 \log_2 n$
 - expected size upper bounded by $n(3/4)^{10 \log_2 n}$
 - $n(3/4)^{10 \log_2 n} \approx n \times n^{-10 \log_2 (4/3)} \approx n^{-3.15}$
 - According Markov's Inequality:

What this mean?

$$\Pr [Y_{10 \log_2 n} \geq 1] \leq E[Y_{10 \log_2 n}] / 1 = n^{-3.15}$$

- the number of steps is $O(\log n)$ with high probability
- Each step has span $O(\log n)$ so the overall span is $O(\log^2 n)$ with high probability



明

SYNOPSIS

- Discrete Probability: Let's Toss Some Dice
- Finding The Two Largest
- Find the k^{th} smallest element
- Quicksort
- Analysis of Quicksort



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Quicksort

- Originally invented and analyzed by Hoare in 1960's
- I strongly urge to watch Jon Bentley on "Three beautiful Quicksorts" at
 - www.youtube.com/watch?v=QvgYAQzg1z8



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Sequential Quicksort



```
int i, j;
for( i = low, j = high - 1; ; )
{
    while( a[ ++i ] < pivot );
    while( pivot < a[ --j ] );
    if( i >= j )
        break;
    swap( a, i, j );
}
// Restore pivot
swap( a, i, high - 1 );
quicksort( a, low, i - 1 ); // Sort small elements
quicksort( a, i + 1, high ); // Sort large elements
```



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Quicksort

- Is there parallelism in quicksort?

```
1  sort  $S$  =  
2  if  $|S| = 0$  then  $S$   
3  else  
4    let  
5       $p = \text{pick a pivot from } S$   
6       $S_1 = \langle s \in S \mid s < p \rangle$   
7       $S_2 = \langle s \in S \mid s = p \rangle$   
8       $S_3 = \langle s \in S \mid s > p \rangle$   
9       $(R_1, R_3) = (\text{sort } S_1 \parallel \text{sort } S_3)$   
10   in  
11      $R_1 ++ S_2 ++ R_3$   
12   end
```

是
創
新




明

Quicksort

- Each call to Quicksort either makes
 - No recursive calls (base case), or
 - Two recursive calls
- **Call tree is a binary**
- Depth the call tree determines the span of the algorithm

<7, 4, 2, 3, 5, 8, 1, 6>



明

Picking The Pivot

- Always pick the first element
 - Worst case $O(n^2)$ work ————— **Why?**
 - In practice, **almost sorted inputs are not uncommon**
- Pick the median of 3 elements (e.g., first, middle and last elements)
 - could possible divide evenly
 - worst case is still bad
- **Pick an element at random**
 - we hope this divides evenly in expectation
 - leading to expected **$O(n \log n)$ work** and **$O(\log^2 n)$ span**



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

SYNOPSIS

- Discrete Probability: Let's Toss Some Dice
- Finding The Two Largest
- Find the k^{th} smallest element
- Quicksort
- Analysis of Quicksort



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Picking The Pivot

- **Pick first element**
 - Worst case $O(n^2)$ work
 - Expected $O(n \log n)$ work
 - ✓ Averaged over all possible orderings
 - Work well on the average
 - Slow on some, possibly common, cases
- **Pick a random element**
 - Expected worst-case $O(n \log n)$ work
 - ✓ For input in any order, the expected work is $O(n \log n)$
 - No input has expected $O(n^2)$ work
 - With a small probability, we could be unlucky and have $O(n^2)$ work



明 Randomized Quicksort

- Assign a uniformly random priority to each number in $[0,1]$

```
1 fun quicksort(S) =  
2   if |S| = 0 then S  
3   else let  
4       val p = pick as pivot the highest priority element from S  
5       val S1 = { s ∈ S | s < p }  
6       val S2 = { s ∈ S | s = p }  
7       val S3 = { s ∈ S | s > p }  
8       val (R1, R3) = (quicksort(S1) || quicksort(S3))  
9   in  
10      append(R1, append(S2, R2))  
11  end
```

- Once the priorities are assigned, the algorithm is deterministic



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明 Randomized Quicksort

- Count comparisons made!
 - Almost all **the work is comparisons**
 X_n = # of comparisons *quicksort* makes on input of size n
- Find $E[X_n]$ for any input sequence S
- Notation:
 - Let **$T = \text{sort}(S)$**
 - T_i and T_j refer to elements in the final sorted order and $i < j$ and **$T_i \leq T_j$**
 - **p_i** refers to priority chosen for T_i
 - **$A_{i,j}$** = 1 if T_i and T_j were ever compared during the sort




華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Analyzing Quicksort

- Crucial point is how to model $A_{i,j}$
- In any one call to quicksort, there are **three cases**
 - Pivot p is either T_i or $T_j \Rightarrow A_{i,j} = 1$
 - $T_i < p < T_j \Rightarrow T_i \in S_1, T_j \in S_2, A_{i,j} = 0$
 - Either $p < T_i$ or $p > T_j \Rightarrow T_i, T_j \in S_1$ or $T_i, T_j \in S_2$
- **If two elements are compared in a quicksort call, they will never be compared again in any other call!**



Why?



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Analyzing Quicksort

$$X_n \leq \sum_{i=1}^n \sum_{j=i+1}^n A_{ij}$$

- The non-optimized code compares each element to pivot 1 times

1 ...

2 $\text{val } S_1 = \langle s \in S \mid s < p \rangle$

3 $\text{val } S_2 = \langle s \in S \mid s = p \rangle$

4 $\text{val } S_3 = \langle s \in S \mid s > p \rangle$

5 ...

- By linearity of expectation

$$\mathbf{E}[X_n] \leq \sum_{i=1}^n \sum_{j=i+1}^n \mathbf{E}[A_{ij}]$$

華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Analyzing Quicksort

- Consider first when the pivot is one of T_i, T_{i+1}, \dots, T_j
- T_i and T_j are compared $\Leftrightarrow p_i$ or p_j is the highest priority among $\{p_i, p_{i+1}, \dots, p_j\}$
 - Assume $T_k, i < k < j$ has higher priority
 - For any subdivision $\dots, T_i, \dots, T_k, \dots, T_j$, T_k will become a pivot and separate T_i and T_j
 - T_i and T_j will never be compared!



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Analyzing Quicksort

$$\begin{aligned} \mathbf{E}[A_{ij}] &= \mathbf{Pr}[A_{ij} = 1] \\ &= \mathbf{Pr}[p_i \text{ or } p_j \text{ is the maximum in } \{p_i, \dots, p_j\}] \\ &= \frac{2}{j-i+1} \text{ (Why ?)} \end{aligned}$$

- $j-i+1$ elements between p_i and p_j and each is equally likely to be the maximum
- We want either p_i or p_j , hence $2/(j-i+1)$
- T_i is compared to T_{i+1} with probability 1



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Analyzing Quicksort

德厚學
求是創新



$$\begin{aligned}\mathbf{E}[X_n] &\leq \sum_{i=1}^{n-1} \sum_{j=i+1}^n \mathbf{E}[A_{ij}] \\ &= \sum_{i=1}^{n-1} \sum_{j=i+1}^n \frac{2}{j-i+1} \\ &= \sum_{i=1}^{n-1} n \sum_{k=2}^{n-i+1} \frac{2}{k} \\ &\leq 2 \sum_{i=1}^{n-1} H_n \\ &= 2nH_n \in O(n \log n)\end{aligned}$$



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

(Recall: $H_n = \ln n + O(1)$.)

明

Analyzing Quicksort

- Indirectly, average work for basic deterministic quicksort is $O(n \log n)$
 - Just shuffle data randomly and apply the basic algorithm
 - \equiv to picking random priorities



Exercise 10.20. Rewrite the quicksort algorithm so that it takes a comparison function $\text{cmp} : \alpha \times \alpha \rightarrow \text{order}$ and a sequence of type α seq , and only uses the comparison once when comparing the pivot with each key. The type order is the set $\{\text{Less}, \text{Equal}, \text{Greater}\}$.

明

Alternative Analysis

- Write a recurrence for the number of comparisons :

$$X(n) = X(Y_n) + X(n - Y_n - 1) + n - 1$$

- Random variable Y_n is the size of S_1

$$\begin{aligned}\mathbf{E}[X(n)] &= \mathbf{E}[X(Y_n) + X(n - Y_n - 1) + n - 1] \\ &= \mathbf{E}[X(Y_n)] + \mathbf{E}[X(n - Y_n - 1)] + n - 1 \\ &= \frac{1}{n} \sum_{i=0}^{n-1} (\mathbf{E}[X(i)] + \mathbf{E}[X(n - i - 1)]) + n - 1\end{aligned}$$

Why?



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Alternative Analysis

$$\begin{aligned}\mathbf{E}[X(n)] &= \frac{1}{n} \sum_{i=0}^{n-1} (\mathbf{E}[X(i)] + \mathbf{E}[X(n-i-1)]) + n - 1 \\ &= \frac{2}{n} \sum_{i=0}^{n-1} \mathbf{E}[X(i)] + n - 1\end{aligned}$$

- With telescoping, this also solves as $O(n \log n)$



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

Expected Span

- S is split into $L(ess), E(qual)$ and $(g)R(eater)$
- Let $X_n = \max\{|L|, |R|\}$
- We use **filter** to partition
 - $S(n) = S(X_n) + O(\log n)$
- The only thing is to calculate the depth of the pivot tree!

Why?



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Expected Span

- the depth of the pivot tree is $O(\log n)$ by relating it to the number of contraction steps of the randomized *kthSmallest*
- In *kthSmallest*, we have analyzed
 - each node had depth greater than $10 \lg n$ with probability at most $1/n^{3.15}$
 - *quicksort* has any node of depth $10 \lg n$ is also $1/n^{3.15}$?
 - It is wrong!



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Expected Span

- we have **multiple nodes** the probability increases that at least one will go above the bound
- Here is where we get to apply the **union bound**
 - For a collection of events A_1, \dots, A_n the bound is

$$\Pr \left[\bigcup_{1 \leq i \leq n} A_i \right] \leq \sum_{i=1}^n \Pr [A_i]$$



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Expected Span

- Here is where we get to apply the **union bound**
 - the **individual events** are the depths of each node being larger $10 \lg n$
 - the **union** is the probability that any of the nodes has depth larger than $10 \lg n$
 - There are **n events** each with probability $1/n^{3.15}$, so the union bound states

$$\Pr [\text{depth of quicksort pivot tree} > 10 \lg n] \leq \frac{n}{n^{3.15}} = \frac{1}{n^{2.15}}$$

Why?



明

Alternative Analysis

- Bernoulli's Inequality

$$(1 + x)^\alpha \geq 1 + \alpha \cdot x \quad (x > -1, \alpha > 0)$$

- Assume

$$x = -\frac{1}{n^{3.15}}, \alpha = n$$

- Then, we can get

$$\left(1 - \frac{1}{n^{3.15}}\right)^n \geq 1 - n \cdot \frac{1}{n^{3.15}} = 1 - \frac{1}{n^{2.15}}$$



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

Alternative Analysis

$$\Pr\left[\sum_{i=1}^n \sum_{j=i+1}^n A_{ij} > 0\right] = 1 - \Pr\left[\sum_{i=1}^n \sum_{j=i+1}^n A_{ij} = 0\right]$$

$$1 - \Pr\left[\prod_{i=1}^n x_i = 0\right]$$

$$= 1 - \prod_{i=1}^n \Pr[x_i = 0]$$

$$= 1 - \left(1 - \frac{1}{n^{3.15}}\right)^n$$

$$\leq 1 - \left(1 - n \cdot \frac{1}{n^{3.15}}\right)$$

$$= \frac{1}{n^{2.15}}$$



华中科技大学

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

明

SUMMARY

- Discrete Probability: Let's Toss Some Dice
- Finding The Two Largest
- Finding The kth Smallest Element
- Quicksort
- Analysis of Quicksort



華中科技大學

HUAZHONG UNIVERSITY OF SCIENCE AND TECHNOLOGY